

Hive

Jay Urbain, PhD

Hive

- SQL-like query language that generates MapReduce
- Developed at Facebook
- It uses H-SQL, not quite the same as ANSI SQL, ~70% overlap
- Underlying execution much different than relational engine
 - MapReduce
- Batch, not interactive.
 - Latency in results
 - Move to Spark, Impala
- Often used with HBase wide column database
- Create schema “on read”

Hbase

- Wide-column, NoSQL database
- Use CREATE TABLE over HDFS
- Once your table is created, query with Hive H-SQL
- Hive libraries are integrated with HBase
- Hive libraries include the H-SQL/HQL language
- Commercial distributions include Hive

Hive usage example:

- `CREATE TABLE table col1, col2`
- `SELECT col1, col2 FROM table`

Hive

Hive SQL Datatypes

INT
TINYINT/SMALLINT/BIGINT
BOOLEAN
FLOAT
DOUBLE
STRING
TIMESTAMP
BINARY
ARRAY, MAP, STRUCT, UNION
DECIMAL
CHAR
VARCHAR
DATE

Hive SQL Semantics

SELECT, LOAD, INSERT from query
Expressions in WHERE and HAVING
GROUP BY, ORDER BY, SORT BY
Sub-queries in FROM clause
CLUSTER BY, DISTRIBUTE BY
ROLLUP and CUBE
UNION
LEFT, RIGHT, and FULL INNER/OUTER JOIN
CROSS JOIN, LEFT SEMI JOIN
Windowing functions (OVER, RANK, etc.)
INTERSECT, EXCEPT, UNION DISTINCT
Subqueries in WHERE
Subqueries in HAVING

	Hive 0.10
	Hive 0.11
	Future

Note: JOINS can have substantial overhead

Working with Hive

- Starting the Hive service
 - Use admin tools to verify its started
 - Turned on by default with commercial distributions
- Managing the metastore database
 - Stores the metadata for Hive tables in a database
 - Can be embedded, local, or remote
 - Schema tool: offline tool for defining schema

Metadata

MySQL	Hive
USE database;	USE database;
SHOW DATABASES;	SHOW DATABASES;
SHOW TABLES;	SHOW TABLES;
DESCRIBE table;	DESCRIBE (FORMATTED EXTENDED) table;
CREATE DATABASE db_name;	CREATE DATABASE db_name;
DROP DATABASE db_name;	DROP DATABASE db_name (CASCADE);

Why use Hive?

- You're an analyst, use databases, spreadsheets
- You know SQL
- Want to ask analytical questions, the questions you ask with SQL
 - Not as useful for data preprocessing
 - Use MapReduce, scripting language, Pig
- Gives convenience of SQL, but is *really* MapReduce.
 - JOINS function differently than on RDBMS
 - NoSQL has no need to be *normalized*, very expensive operation on big data

WordCount in Hive

- Pseudo code:
 1. Create HBase table for text input
 2. Load your data into the table
 3. Create a new destination table
 4. Load the table with processed results

WordCount in Hive

```
CREATE TABLE wordcount AS
SELECT word, count(1) AS count
FROM (SELECT EXplode(SPLIT(LCASE
    (REGEXP_REPLACE
        (line, '[\p{Punct},\p{Cntrl}]', '')), ' '))
    AS word FROM myinput) words
GROUP BY word
ORDER BY count DESC, word ASC;
```

Regex is ugly, best to use something else for preprocessing

Using Hive

- HBase tables and Hive queries
- Query optimization for Hive
- Partitioning, bucketizing, or sampling (subsets)
- Cost-based optimization (CBO)
- Column based statistics

Reading HQL Query Plans

- Read from bottom to top:

```
[impalad-host:21000] > explain select count(*) from customer_address;
+-----+
| Explain String                                     |
+-----+
| Estimated Per-Host Requirements: Memory=42.00MB VCores=1 |
|                                                         |
| 03:AGGREGATE [MERGE FINALIZE]                         |
| | output: sum(count())                               |
| |                                                     |
| 02:EXCHANGE [PARTITION=UNPARTITIONED]                 |
| |                                                     |
| 01:AGGREGATE                                           |
| | output: count(*)                                   |
| |                                                     |
| 00:SCAN HDFS [default.customer_address]              |
| partitions=1/1 size=5.25MB                           |
+-----+
```

Hue Home -> examples -> Salary Growth (Hive Query)

Hue - Welcome Home - Mozilla Firefox

about:sessionrestore x Hue - Workflow Editor x Hue - Welcome Home x

quickstart.cloudera:8888/home?uuid=ccd89375-5125-4d1a-a278-8a7c58f977c7

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

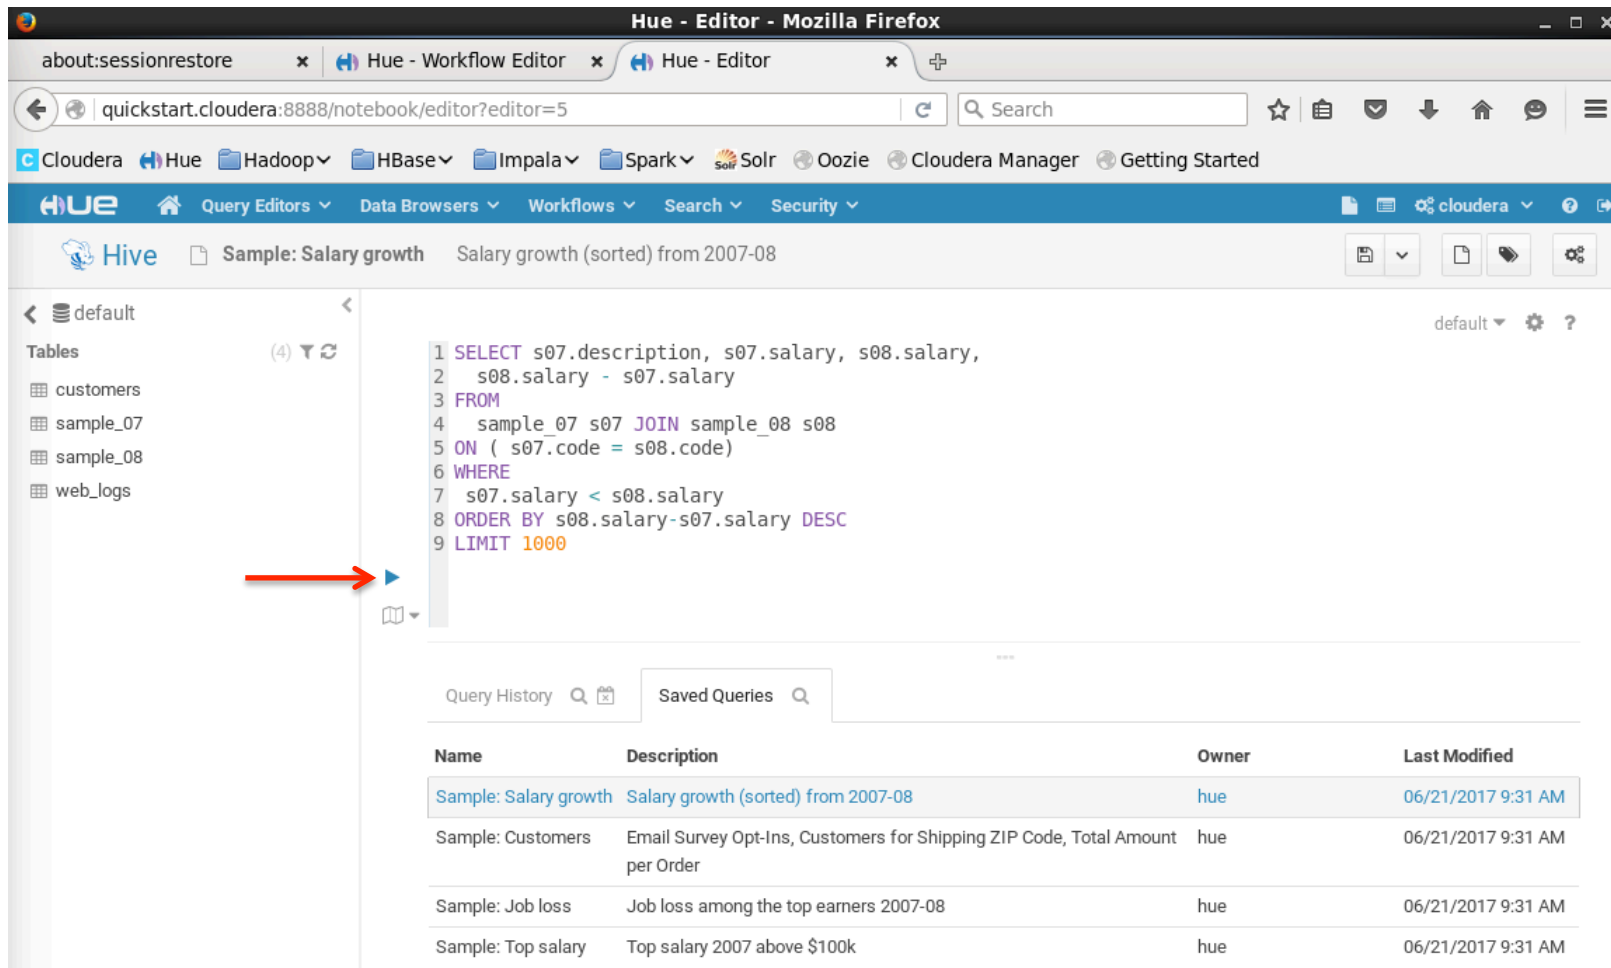
HUE Query Editors Data Browsers Workflows Search Security

My documents

My documents > examples

Name	Type	Owner	Last Modified
Sample: Customers	Hive Query	hue	06/21/2017 9:31 AM
Sample: Customers	Impala Query	hue	06/20/2017 7:33 AM
Sample: Job loss	Hive Query	hue	06/21/2017 9:31 AM
Sample: Job loss	Impala Query	hue	06/20/2017 7:33 AM
Sample: Salary growth	Hive Query	hue	06/21/2017 9:31 AM
Sample: Salary growth	Impala Query	hue	06/20/2017 7:33 AM
Sample: Top salary	Hive Query	hue	06/21/2017 9:31 AM

Execute Query



Hue - Editor - Mozilla Firefox

about:sessionrestore x Hue - Workflow Editor x Hue - Editor x

quickstart.cloudera:8888/notebook/editor?editor=5

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

HUE Query Editors Data Browsers Workflows Search Security

Hive Sample: Salary growth Salary growth (sorted) from 2007-08

default

```
1 SELECT s07.description, s07.salary, s08.salary,
2       s08.salary - s07.salary
3 FROM
4       sample_07 s07 JOIN sample_08 s08
5 ON ( s07.code = s08.code)
6 WHERE
7       s07.salary < s08.salary
8 ORDER BY s08.salary-s07.salary DESC
9 LIMIT 1000
```

Query History Saved Queries

Name	Description	Owner	Last Modified
Sample: Salary growth	Salary growth (sorted) from 2007-08	hue	06/21/2017 9:31 AM
Sample: Customers	Email Survey Opt-Ins, Customers for Shipping ZIP Code, Total Amount per Order	hue	06/21/2017 9:31 AM
Sample: Job loss	Job loss among the top earners 2007-08	hue	06/21/2017 9:31 AM
Sample: Top salary	Top salary 2007 above \$100k	hue	06/21/2017 9:31 AM

Hue - Editor - Mozilla Firefox

about:sessionrestore

Hue - Workflow Editor

Hue - Editor

quickstart.cloudera:8888/notebook/editor?editor=50020

Search

Cloudera

Hue

Hadoop

HBase

Impala

Spark

Solr

Oozie

Cloudera Manager

Getting Started

HUE

Query Editors

Data Browsers

Workflows

Search

Security

cloudera

Hive

Sample: Salary growth

Salary growth (sorted) from 2007-08

default

Tables

customers

sample_07

sample_08

web_logs

55.117s

default

1

2

3

4

5

6

7

8

9

SELECT s07.description, s07.salary, s08.salary,

s08.salary - s07.salary

FROM

sample_07 s07 JOIN sample_08 s08

ON (s07.code = s08.code)

WHERE

s07.salary < s08.salary

ORDER BY s08.salary-s07.salary DESC

LIMIT 1000

Query History

Saved Queries

Results (767)

s07.description

s07.salary

s08.salary

_c3

1 Dentists, all other specialists

120360

142070

21710

2 Surgeons

191410

206770

15360

3 Oral and maxillofacial surgeons

178440

190420

11980

4 Natural sciences managers

113170

123140

9970

5 Physicians and surgeons, all other

155150

165000

9850

6 Orthodontists

185340

194930

9590

7 Internists, general

167270

176740

9470

8 Political scientists

90050

99320

9270

9 Obstetricians and gynecologists

183600

192780

9180

10 Chief executives

151370

160440

9070

Explain

The screenshot shows the Hue - Editor interface in Mozilla Firefox. The browser address bar displays `quickstart.cloudera:8888/notebook/editor?editor=50020`. The interface includes a top navigation bar with links to Cloudera, Hue, Hadoop, HBase, Impala, Spark, Solr, Oozie, Cloudera Manager, and Getting Started. Below this is a blue navigation bar with tabs for Query Editors, Data Browsers, Workflows, Search, and Security. The main workspace is titled "Sample: Salary growth" and "Salary growth (sorted) from 2007-08". On the left, a sidebar shows a file tree with "default" and "Tables" (customers, sample_07, sample_08, web_logs). The central area contains a SQL query editor with the following code:

```
1 SELECT s07.description, s07.salary, s08.salary,  
2       s08.salary - s07.salary  
3 FROM  
4   sample_07 s07 JOIN sample_08 s08  
5   ON ( s07.code = s08.code)  
6 WHERE  
7   s07.salary < s08.salary  
8 ORDER BY s08.salary-s07.salary DESC  
9 LIMIT 1000
```

A red arrow points to the "Explain" button in the toolbar, which is located below the query editor. The toolbar also includes buttons for "Format" and "Clear". Below the query editor, the results are displayed in a table with the following columns: "description", "s07.salary", "s08.salary", and "_c3". The table contains 10 rows of data, sorted by the difference in salary between s08 and s07 in descending order.

	description	s07.salary	s08.salary	_c3
1	Dentists, all other specialists	120360	142070	21710
2	Surgeons	191410	206770	15360
3	Oral and maxillofacial surgeons	178440	190420	11980
4	Natural sciences managers	113170	123140	9970
5	Physicians and surgeons, all other	155150	165000	9850
6	Orthodontists	185340	194930	9590
7	Internists, general	167270	176740	9470
8	Political scientists	90050	99320	9270
9	Obstetricians and gynecologists	183600	192780	9180
10	Chief executives	151370	160440	9070

Hue - Editor - Mozilla Firefox

about:sessionrestore x Hue - Workflow Editor x Hue - Editor x

quickstart.cloudera:8888/notebook/editor?editor=50020

Search

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

HUE Query Editors Data Browsers Workflows Search Security cloudera

Hive Sample: Salary growth Salary growth (sorted) from 2007-08

default

Tables (4)

customers

sample_07

sample_08

web_logs

```
1 SELECT s07.description, s07.salary, s08.salary,
2       s08.salary - s07.salary
3 FROM
4       sample_07 s07 JOIN sample_08 s08
5 ON ( s07.code = s08.code)
6 WHERE
7       s07.salary < s08.salary
8 ORDER BY s08.salary-s07.salary DESC
9 LIMIT 1000
```

Query History Saved Queries Results (767) Explain

STAGE DEPENDENCIES:

Stage-5 is a root stage

Stage-2 depends on stages: Stage-5

Stage-0 depends on stages: Stage-2

STAGE PLANS:

Stage: Stage-5

Map Reduce Local Work

Alias -> Map Local Tables:

s07

Fetch Operator

limit: -1

Alias -> Map Local Operator Tree:

s07

TableScan

alias: s07

Statistics: Num rows: 225 Data size: 46055 Basic stats: COMPLETE Column stats: NONE

Filter Operator

predicate: code is not null (type: boolean)

Statistics: Num rows: 113 Data size: 23129 Basic stats: COMPLETE Column stats: NONE

HashTable Sink Operator

keys:

0 code (type: string)

1 code (type: string)

Review

- All are true about Hive, except ____
 - It is SQL-like query
 - It uses H-SQL
 - It is used in conjunction with HBase
 - It is interactive
- Which of the following data types was supported in Hive 0.11
 - Varchar
 - Binary
 - Float
 - Decimal
- People who are mostly Analysts who know SQL, and who want to ask analytical questions will most likely use Hive over other methods
 - False
 - True
- Some of the query optimization for Hive includes which of the following?
All of these answers
 - Partitioning, bucketizing, or sampling
 - Using column-based statistics
 - Cost-base optimization