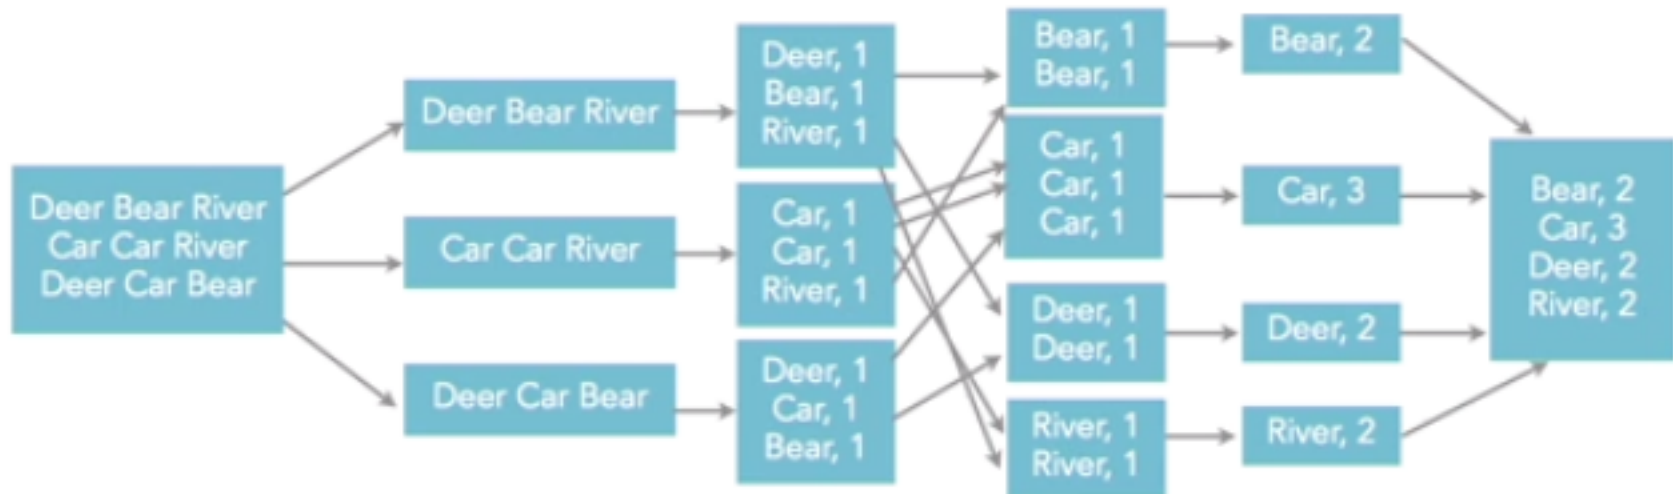


Tuning MapReduce

Jay Urbain, PhD

Tuning by physical methods

Input Splitting Mapping Shuffling Reducing Final Result



Better MapReduce: Optimizations

- Optimize before job runs
 - Preprocessing on data, files size, compression, encryption
- Optimize onload of the data
 - Change compression ratio, different size chunks
- Optimize the map phase of the job
 - Code, hashing
- Optimize shuffle phase of the job
 - Get better distribution on shuffle-sort
- Optimize the reduce phase
 - Additional transformation, combine or separate aggregations
- Optimize after job completes
 - Partition job in to separate jobs, pipeline
- *Other*
 - *Make sure you verify your cluster parameters especially for nonstandard configurations*
 - *Unused resources*
 - *Overstressed resources – spill out to disk – use cluster-monitoring tools, e.g., Ganglia*
 - *Add more nodes – especially when you're using cloud*
 - *Mixed local and cloud Hadoop*

Understanding MapReduce

- Map
 - $(k1, v1)$: input information
 - $(k2, v2)$: intermediate values (one list per node)
 - can use *combiner* (mini-Reducer in Mapper) before emitting intermediate values
- Shuffle/sort
- Reduce
 - $(k2, \text{list}(v2))$
 - $(k3, v3)$
- Notes:
 - Keep jobs as simple as possible, consider adding steps
 - One type of activity at a time

Mapper Task Optimization

- Sub-divide tasks
 - Chain jobs
 - 1-3 minutes per map task run
- Custom partitioner
 - Default is hash partitioner
- Reduce the amount of data
 - skip unnecessary or bad data, add counter, partition input files
- Logging and counters
- Monitor/tune for optimal spill ratio
- Local Reducers/ combiners
- Set Map-only jobs
 - Photo processing – no aggregation needed

Input file compression

- Vendor implementations differ
- Read the documentation
- Cloudera's version of LZO is 'splittable'
- Use best from vendor for app

Data Types

- Custom types supported – write `RawComparator`
 - Custom comparator
 - E.g. geospatial data
 - Custom input split
 - Custom partitioner
 - Use the Text type only when you are using text data

Reducer Task Optimization

- Subdividing tasks (chaining jobs)
- Logging/debugging to see what's going on
- Secondary sort
- Set thresholds
- May not need a Reducer (Map only)
 - Photos, music, transformation only apps
- Reduce data in Mapper

`org.apache.hadoop.mapreduce.filecache.DistributedCache`

Creating a Distributed Cache

- Used to cache files at nodes needed to run application jobs
- Include the path via URL's (hdfs:// or http://) in JobConf settings
- Files will be copied to each data node on job execution
- E.g.:
 - Using a lookup file for translation, black list, white list
 - Translate numeric code to human values

Summary

- Add more node
 - Especially in cloud environment
- Distributed cache
- Skipping bad records
- Subdividing tasks
- Logging/profiling
- Debugging, unit testing
- Optimized Java processing
- String manipulation can be a bottleneck
 - `StringBuffer.append`
 - `LongWritable` and `BytesWritable` are faster than text parsing

Review

- Which of the following is a Mapper Task Optimization?
 - Subdivide tasks
 - All of the answers
 - Reduce the amount of data
 - Monitor/tune for optimal spill ratio
- Creating a distributed cache is important for optimization; it is used to cache files like text files, archives, and JAR files that are needed to run applications or jobs
 - True
 - False

Review

- Reducer task optimization usually falls into four categories: subdivide tasks, logging/debugging, secondary sort, and setting thresholds
 - True
 - False